Benjamin Mah

+1 647-901-5309 | benjaminmah.bm@gmail.com | benjaminmah.com | linkedin.com/in/benjaminmahh | github.com/benjaminmah

Education

University of Toronto June 2026

Bachelor of Applied Science in Engineering Science (Honours) *Toronto, Canada*

Major in Machine Learning Engineering, Certificate in Engineering Business

Relevant Courses: Data Structures & Algorithms, Neural Networks, Deep Learning, Distributed Systems, Probabilistic Reasoning

Experience

Asana August 2026 – Present

Data Scientist *Vancouver, Canada*

- Designed, ran, and analyzed company-level A/B tests optimizing conversion and retention on self-serve, product-led growth (PLG) experiences across **200,000+** customers, driving **$3M+ ARR** in shipped wins, owning the ship/no-ship decision based on causal impact analysis, and driving ideation for iterative follow-up experiments
- Served as the data owner for the AI-powered Meetings feature, the flagship feature of the company's expansion from **1 product to 5**, defining success and adoption metrics and reporting performance to product leadership
- Partnered with PMs, engineers, and designers across multiple teams on **20+** ad hoc analysis requests, building **Databricks SQL** dashboards to track AI feature adoption, usage, and consumption-based monetization
- Supported UX research by pulling targeted cohorts of **1,000+** users to gather input on new features via surveys

Asana May 2025 – August 2025

Data Science Intern *Vancouver, Canada*

- Led data analysis of key patterns in customer churn using **SciPy** for clustering and **XGBoost** for predictive modeling, uncovering a potential **$2M+ ARR** recovery through proactive, personalized retention strategies
- Designed and executed A/B tests for self-serve product feature changes, analyzing business impact metrics across **40,000+ unique users** to inform data-driven shipping decisions that increased conversion by **$1.2M ARR**
- Conducted opportunity sizing analysis for automatic seat upgrades using **Databricks** and **Redash**, examining upgrade-to-churn patterns across **7,000+ organizations** monthly to prioritize ARR optimization opportunities

Mozilla May 2024 – April 2025

Machine Learning Engineer Intern *Toronto, Canada*

- Built and deployed ML models using **XGBoost** and **scikit-learn** for automated bug triaging through a bot deployed in **Bugzilla** to streamline triage workflows and reduce manual processing by **50+ hours** monthly
- Fine-tuned LLMs using **Hugging Face Transformers**, **PyTorch**, and **PEFT** on **15,000+** patch fixes to generate code addressing review comments, achieving an **8%** higher acceptance rate than **GPT-4** and **Claude 3**
- Developed a code review bot used by **200+ engineers**, reducing review time by **30%** on production pull requests

Cohere September 2023 – April 2024

Data Specialist *Toronto, Canada*

- Audited and annotated machine learning datasets using **Scale AI**, identifying and correcting labeling errors across thousands of samples, enhancing model performance and saving over **100 hours** of processing time each month

Royal Bank of Canada May 2023 – August 2023

Machine Learning Engineer Intern *Toronto, Canada*

- Built a semantic search tool for financial advisors using **LangChain**, **Qdrant**, and **Dask**, boosting information retrieval and extractive question answering speed by **40%** through optimized query embeddings and indexing
- Designed a scalable multi-format file processing pipeline with retrieval endpoints using **Python** and **Amazon S3**, cutting proof-of-concept development time by **10+ hours** per project and standardizing reusable code

Publications

Dissecting Subjectivity and the "Ground Truth" Illusion in Data Annotation | *ACM FAccT 2026*

Impact of an LLM-based Review Assistant in Practice | *IEEE Transactions on Software Engineering 2026*

Skills

Languages: Python, SQL, C/C++, JavaScript, HTML/CSS, MATLAB
Frameworks/Tools: PyTorch, TensorFlow, Databricks, Redash, LangChain, AWS, Flask, Git, VSCode, Jupyter
Libraries/APIs: pandas, NumPy, scikit-learn, XGBoost, Transformers, SciPy, Matplotlib, ChromaDB, Qdrant