

Benjamin Mah

+1 647-901-5309 | benjaminmah.bm@gmail.com | benjaminmah.com | linkedin.com/in/benjaminmahh | github.com/benjaminmah

EDUCATION

University of Toronto

June 2026

Bachelor of Applied Science in Engineering Science (Honours)

Toronto, Canada

Major in Machine Learning Engineering, Certificate in Engineering Business

Relevant Courses: Data Structures & Algorithms, Neural Networks, Deep Learning, Distributed Systems, Probabilistic Reasoning

EXPERIENCE

Asana

August 2026 – Present

Data Scientist

Vancouver, Canada

Asana

May 2025 – August 2025

Data Science Intern

Vancouver, Canada

- Led comprehensive data analysis to identify key patterns in customer churn using clustering and predictive churn modeling, revealing a potential recovery of **\$2M+ ARR** through proactive and personalized retention strategies
- Designed and executed A/B tests for self-serve product feature changes, analyzing business impact metrics across **40,000+ unique users** to inform data-driven shipping decisions that increased conversion by **\$1.2M ARR**
- Conducted opportunity sizing analysis for automatic seat upgrades using **Databricks** and **Redash**, examining upgrade-to-churn patterns across **7,000+ organizations** monthly to prioritize ARR optimization opportunities

Mozilla

May 2024 – April 2025

Machine Learning Engineer Intern

Toronto, Canada

- Built and deployed ML models using **XGBoost** and **scikit-learn** for automated bug triaging through a bot deployed in **Bugzilla** to automate triage workflows and reduce manual processing by **50+ hours** monthly
- Co-authored a research paper on a GPT-based code review bot, collaborating with **Ubisoft**, **Queen's University**, and **Polytechnique Montréal**, used by over **200 engineers** and leading to a **30%** reduction in review time
- Fine-tuned pre-trained large language models using **Hugging Face Transformers**, **PyTorch**, and **PEFT** to generate code addressing inline comments, achieving an **8%** higher acceptance rate with a dataset of **15,000+** annotated patch fixes and review comments compared to baseline models like **GPT-4** and **Claude 3**

Cohere

September 2023 – April 2024

Data Specialist

Toronto, Canada

- Audited and annotated machine learning datasets using **Scale AI**, identifying and correcting labeling errors across thousands of samples, enhancing model performance and saving over **100 hours** of processing time each month
- Conducted in-depth comparative analysis of LLM-generated code completions written in **Python**, **JavaScript**, and **SQL**, refining validation processes to align with coding standards, resulting in faster deployment cycles

Royal Bank of Canada

May 2023 – August 2023

Machine Learning Engineer Intern

Toronto, Canada

- Developed a machine learning tool for financial advisors, using semantic search with **LangChain**, **Qdrant**, **Dask**, and **NumPy** to efficiently process user queries from financial databases, boosting response speed in information retrieval (IR) and extractive question answering (EQA) by **40%** through optimized query embeddings
- Designed and implemented a scalable multi-type file processing pipeline with efficient retrieval endpoints, using **Python** and **Amazon S3** to handle various complex textual data formats, reducing proof-of-concept (PoC) development time by **10 hours** per project and standardizing reusable code for future generative AI initiatives

PUBLICATIONS

Dissecting Subjectivity and the “Ground Truth” Illusion in Data Annotation  | *ACM FAccT 2026*

Impact of an LLM-based Review Assistant in Practice  | *IEEE Transactions on Software Engineering 2026*

SKILLS

Languages: Python, SQL, C/C++, JavaScript, HTML/CSS, MATLAB

Frameworks/Tools: PyTorch, TensorFlow, Databricks, Redash, LangChain, AWS, Flask, Git, VSCode, Jupyter

Libraries/APIs: pandas, NumPy, scikit-learn, XGBoost, Transformers, SciPy, Matplotlib, ChromaDB, Qdrant